

Learning Objectives

The learning objectives in the tutorial are to understand the individual steps of the *Secure Workflow*. This should provide a deeper understanding of the individual measures taken to utilize secure compute capabilities on a shared HPC system.

Warning: Usually this Secure Workflow would be automated End-to-End on a Secure Client and the user would only call it in a parametrized fashion, e.g. through a GUI. In this tutorial we will manually go through the steps to get a detailed understanding.

Contents

The Secure Client 1: Tutorial (10 min)	1
Encrypting the Data 2: Tutorial (10 min)	2
Building an Encrypted Singularity/Apptainer Container 3: Tutorial (10 min)	3
Uploading Keys and Writing a Batch Script 4: Tutorial (10 min)	3
Encrypting your Batch Script 5: Tutorial (10 min)	4
Signing your Batch Script and Securely Submitting Your Job 6: Tutorial (10 min)	5
Optional: Quick Summary if You Need to Repeat 7: Tutorial (10 min)	5
Optional: Getting to Know the Job 8: Tutorial (10 min)	6

This Tutorial will guide you step-by-step through the *Secure Workflow*. It was developed to provide secure processing capabilities for sensitive data, e.g. medical data. In order to allow this, it needs to be ensured that **only** the legitimate user can access this data, explicitly excluding admins and, of course, attackers. The secure workflow is sketched in fig. 3.

The fundamental idea is that from a secure client system, which could be a system located in a secured hospital infrastructure, all data is encrypted and uploaded to the (potentially unsecure) shared HPC system. The keys **must not** be uploaded to the HPC system, since this would render the encryption useless, because an attacker would then get access to both, your data and your encryption key. Instead we are using a separate key management system, which is outside of the HPC system, and therefor shielded from an attacker on the HPC system. Since the data need to be decrypted during computations, dedicated and isolated compute nodes are required. These compute nodes are part of a dedicated, hidden partition **secure**.

The Secure Client 1: Tutorial (10 min)

As described in the synopsis, the secure client is always the beginning of a secure workflow. You should have gotten a uid. Since we will need to connect from the secure client machine to the HPC system, we need to activate **AgentForwarding**. This can be done with the **-A** option, or using a local config file in `/.ssh/config`.

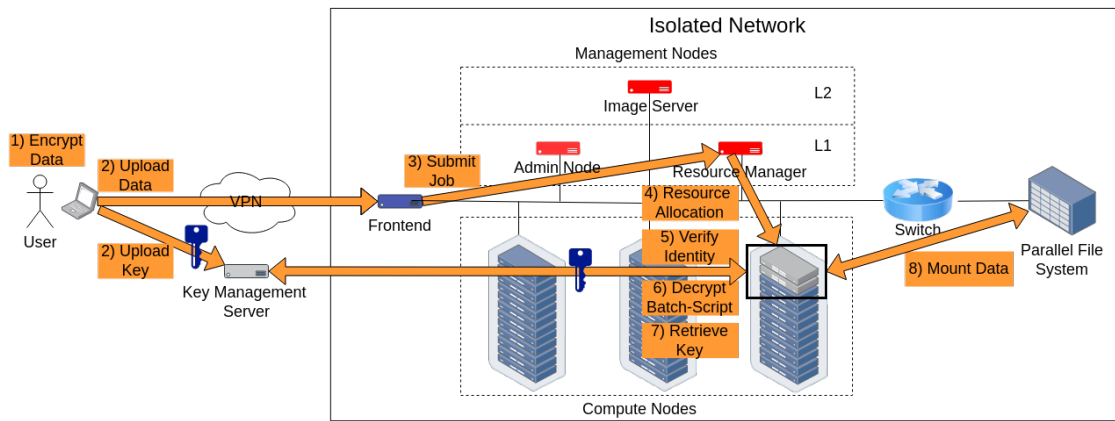


Figure 1: A schematic sketch of the secure workflow on an HPC system, which is divided into 8 distinct steps. As shown, the sensitive data along with the container and the batch script are encrypted on the local machine of the user. Additionally, the batch script is also signed by the user. After encrypting, the components are safe to be uploaded into the shared file system of the HPC system. Although the batch script is signed and encrypted, it can be normally submitted in the third step. As usual, the resource manager will allocate the required resources and start the job in step 4). In the fifth step, the authenticity of the user request is verified and in the 6th step, the batch script is decrypted. Using the provided token for the key management system (KMS) the keys are retrieved in step 7 and used to decrypt the data and container on the isolated, secure node in step 8.

Please connect now to the secure client with:

```
$ ssh -A <uid>@141.5.99.62
```

Or prepare your `/.ssh/config` file like this:

Listing 1: `/.ssh/config`

```
1 Host secure_tutorial
2   User <uid>
3   Hostname 141.5.99.62
4   IdentityFile ~/.ssh/<your_priv_key>
5   ForwardAgent yes
```

And then do a

```
$ ssh secure_tutorial
```

You probably need to choose your favorite shell, e.g. type:

```
$ bash
```

and you will be in a Bash-Shell.

You should now be in your `$ $HOME` and you can enter our `job_template` directory. Here you will find `src` directory containing a Singularity/Apptainer recipe. This can be used to build your on container image to run *BART*. The `data` directory contains your sample input data. In a real use case, this would be your medical data from a patient, which requires the highest protection.

Encrypting the Data 2: Tutorial (10 min)

The first step in the secure workflow is to encrypt your data. In this example, we will be using LUKS. In order to create a LUKS data container you can use a helper script located in `/opt/secure_workflow/create_data_container.s`. Executing it with the `-h` option provides you with information on how to execute it. It will show you, that you need to provide 3 parameters:

- The container name, which is an arbitrary name you can choose, e.g. myinputdata.
- The mount path on your local system, e.g. /mnt/<uid>, or home/<client-uid>/<mymountpoint>
- The size of your LUKS data container in MB to the basis of 10

So please go ahead on type:

```
$ /opt/secure_workflow/create_data_container.sh <containername> /mnt/<uid> 50
```

Please notice, that you need to provide a minimal size of 50MB.

You can verify, that everything worked properly by doing a `$ ls /mnt/<uid>/<containername>`. If you see a fresh ext4 filesystem, indicated by the `lost+found` directory, it worked. You can now copy your input data into that container. Everything you write into that container is automatically encrypted. You can use for instance:

```
$ cp data/* /mnt/<uid>/<containername>
```

```
$ ls /mnt/<uid>/<containername>
```

You can now continue to umount the LUKS data container in order to be able to safely upload it to the HPC-system. We also have an wrapper script for that, you can just do:

```
$ /opt/secure_workflow/umount_data_container.sh <containername> /mnt/<uid>
```

If you look into your `$pwd` you will see two new files:

- <containername>.img
- <containername>.key

The *.img is your LUKS data container, which contains your encrypted data. This is secure to upload to the shared HPC system. The *.key file contains the key necessary to decrypt your data. This **must never** be uploaded to the HPC system, since an attacker would then have access to the encrypted data and also to the decryption key. Instead we will upload it to the dedicated Key Management System. But first we upload the LUKS data container:

Please ensure beforehand that the directory exists. If it doesn't, you can do:

```
$ ssh <hpc-uid>@gwdu101.gwdg.de 'mkdir /scratch/users/<hpc-uid>/secure'
```

Then you can go ahead and upload your LUKS container into this directory:

```
$ scp <containername>.img <hpc-uid>@transfer-scc.gwdg.de:/scratch/users/<hpc-uid>/secure
```

Please notice that the LUKS container has to be in exactly this folder.

Building an Encrypted Singularity/Apptainer Container 3: Tutorial (10 min)

In order to build a Singularity Container, you can use a helper script `buildEncryptedSingularity.sh`. Here, we use an RSA keypair to build an encrypted Singularity/Apptainer image with the recipe which you can find in `src/bart_fft.def`. You can do the following:

```
$ ./buildEncryptedSingularity.sh
```

```
$ ls -lah bart_fft_enc.sif
```

Don't be confused if you see an Error message. One test might fail.

You should see, that you successfully build your encrypted container. In order to be able to run it, you need your private key `rsa_pri.key`. Since the container is encrypted you can securely upload it:

```
$ ssh <hpc-uid>@gwdu101.gwdg.de 'mkdir /scratch/users/<hpc-uid>/home'
```

```
$ scp bart_fft_enc.sif <hpc-uid>@transfer-scc.gwdg.de:/scratch/users/<hpc-uid>/home
```

Uploading Keys and Writing a Batch Script 4: Tutorial (10 min)

Due to the encryption of our data and the encryption of the Singularity/Apptainer images, we have two keys, which we need to manage. These keys are uploaded into our Key Management System, i.e. Vault. In order to retrieve the keys later, we are receiving single-use and short-lived tokens from Vault. These tokens need to be

included in the batch script, so that the keys can be retrieved on a secure node.

The general *command*, i.e. the bash-commands, is written in `command.sh.template`. Please open it in an editor, like vim, emacs or nano and substitute **only** the parameterized variables `<hpc-uid>`, `<uid>`, `<LUKScontainername>` and `<containername>` in **both** scripts. (Please Notice that unfortunately `<LUKScontainername>` is the same as `<containername>`) Please do the same at the bottom of `prepare_scripts.sh`.

Afterwards you can look at `prepare_scripts.sh` and substitute the key name for your LUKS container. You can then execute it with:

```
$ ./prepare_scripts.sh
```

This will upload your keys to Vault, and will insert the received tokens into a new file called `command.sh`. This was automatically derived by `prepare_scripts.sh` by your `command.sh.template`.

To summarize, your resulted `command.sh` is the bash-script, that you want to execute on a secure node.

Encrypting your Batch Script 5: Tutorial (10 min)

Since it does contain the tokens to retrieve the keys from Vault, we should encrypt it as well. This can be done with `gpg`, for which a public key is already imported on the secure client. In addition, there is a helper script in `/opt/secure_workflow/encrypt_script.sh` which you can call with:

```
$ /opt/secure_workflow/encrypt_script.sh command.sh agq001
```

If you get a warning, no problem, just go ahead and say yes

This will encrypt your `command.sh` and will store it in `command.sh.asc`.

If you have a look at it, e.g. with `cat`, you will see that this is no bash script anymore, but an encrypted `gpg` message. Since a valid bash script is required by Slurm, we need to provide one. For this passing in the encrypted `gpg` message to a `decrypt_and_execute` function, which we provide you on a secure node. The script that we want to submit to Slurm is called `run.sh`. Therefor, an encrypted script has to look exactly like this(of course, the actual message will be different):

Listing 2: `run.sh`

```
1  #!/bin/bash
2
3  /usr/bin/decrypt_and_execute <<EOF
4  -----BEGIN PGP MESSAGE-----
5
6  hQIMA8ErHWKPrkmoAQ//daYFx4qwwc72XY2WavHv4RiEEArHqUXL1jVCATi/h8+c
7  hPkKL65KiAvExQgarvgJCRDhLt47F1HSUR3NT4pq1n1/Ql10r2eSksUx8G60o0ia
8  KIDBceAyTe9joEb9PYMXHe4K2pJBt17GIWyBlX0c4ylzh/srSgbTNgxz7GnYJ8Qs
9  KwwjEpFa0y8f4Fuwir9ouKqNZnuQLr/df6WDHAZYSQwvJcn26XhudvEEmQt613x
10 L/13fdWytwKoKvoqent+VuvXtYnq84Sk/rab+HwzuxfVuVgUEYelyBQrOpyy62Bj
11 fQ1ZGWNigjjfQXY+pxHhDuk08CIMWUBfVpZ0To3D/ppnJiJvIjhRu9zvM3vZYXLt
12 7Nb00ZPQDdiaDoap1k7msPh49QmdPk9VNgEPLx/fmMW9NJYXLG0cMv9JDTX7zguoT
13 eJEBmyZoo7mL5XPycPCPoYZrZcVvYjCgya6uWZp6tOYCuwIQzf7qM+puqlS3WnEg
14 rZvgiD1Ju80B6pB5cHPDhuB8+x7M9vSSzFF8LwuCj+/YfYPQkymoGdSaDIYdB2Fz
15 caqAhlWd/Z8MG91+8JZbUpuLkg57jTWqI7qSSmHPJZu7rpAe5p8nTHRMhZ8Idki0
16 uk06zSoQX42PaUPZMQZfAJE/NNxC5uAJ8xkBMzXF/ymRaboyETuxalmw3vyX1J7S
17 wTcBi66cufsqH45W4srKptKNKQpAf5M9JfRmXfHsQBzxwh8mWJ9gFmsrsaC6tTK1
18 KLeLapzVafnIk8/nvjxBmKRSc35GosWTi/AFKk8gj4FuSYtiUbFk8qggcXEVf18v
19 rEOVHAUIkFcdJ78y0slEx4ARu3vkC0+8+ublmieF61zRIh6ETdQAVpbrSfEVUnb
20 F/MJHeIrbED4tFp7VpwRw0459b4E77sMSAOI1xbWXWcpJNLrkR3K5fCsYrxnYd
21 r22ew3T3gAGX2SzaAYgaWGLHeLODvQMIgwBT47Fe320jbngXMD1AAz5YqjKkLxMN
22 onErff2US19R5RXDAHOpAMTQ+u4XVGCx6rBE/4ySV44tbV+6qyy68/dBzP0o+Zvz
23 Jnp5nmlpu8h3u6MTmOKypwQUT4uxfe5J04NtIHT2GzYm1yINcOBBJqb3+MI2SNJ
24 S5us8whI/Srwz1J0PWCpx3XR/IoVe+rrw560ik2NaGDd0rk9IhHLbd+NAXm2aC9P
25 9BJbdyqiU7TxyZ7+2WKycTYULz67qo/B+On2YjWwh24yYWKbFYZPjCq1vhyfCNPa
26 b06EDMdhYcTkFFOJABd/Lx4v8df035hFo11EmawYG7bIvF4qEFDRPLR0CrbII89h
27 /02S+G/v2huFNK1z095WKKXk0Pqb1GiNlvg==
28 =Vq+R
29 -----END PGP MESSAGE-----
30 EOF
```

Please notice the peculiar looking two EOF's!
Of course, you need to insert your own gpg message

Signing your Batch Script and Securely Submitting Your Job 6: Tutorial (10 min)

Since the integrity of the secure nodes depends on the restricted access, to prevent an attacker from accessing the node and tampering with it, we need to sign our batch script, to authenticate ourselves on the secured node. For this, a public-private key pair was generated, where the public key is imported on the secure node and the private key is only stored on the secure client. With this private key, we can do a detached signature.

Please realize: The integrity of this entire workflow depends on the security of this private key with which we authenticate ourselves.

You can perform the detached signature with:

```
$ gpg --detach-sign --local-user user_key -o run.sh.sig run.sh
```

You will be prompted for a passphrase. It is: test

This will read in your `run.sh` create an additional file called `run.sh.sig`. You will need both files to submit your job on the HPC-System. Please upload both of them into the same folder:

```
$ scp run.sh run.sh.sig <hpc-uid>@gwdu101.gwdg.de:/scratch/users/<hpc-uid>/home/
```

You can now submit your job with:

```
$ ssh <hpc-uid>@gwdu101.gwdg.de 'cd /scratch/users/<hpc-uid>/home ; /opt/slurm/bin/sbatch -p secure --exclusive run.sh'
```

After a few seconds, your job will be finished.

The result should be stored in the same LUKS data container. You can fetch it again via `scp` as follows:

```
$ scp <hpc-uid>@transfer-scc.gwdg.de:/scratch/users/<hpc-uid>/secure/<containername>.img .
```

You can then mount the container (provided that your key lies in the same directory) with:

```
$ /opt/secure_workflow/mount_container.sh <containername>
```

and then you can:

```
$ ls -lah /mnt/<uid>/<containername>/
```

You should see two new files: `output_fft.cfl` and `output_fft.hdr`

Optional: Quick Summary if You Need to Repeat 7: Tutorial (10 min)

This is a quick summary if you encountered at some point an Error.

Create an LUKS data container:

```
$ /opt/secure_workflow/create_data_container.sh <containername> /mnt/<uid> 50
```

Copy your data into it:

```
$ cp data/* /mnt/<uid>/<containername>
```

```
$ ls /mnt/<uid>/<containername>
```

And umount it: `$ /opt/secure_workflow/umount_data_container.sh <containername> /mnt/<uid>`

Make all the necessary changes in `prepare_script.sh` and `command.sh.template` and then execute:

```
$ ./prepare_scripts.sh
```

Encrypt your resulting `command.sh` which contains your KMS tokens:

```
$ /opt/secure_workflow/encrypt_script.sh command.sh agq001
```

Make a detached signature to authenticate on the secured node:

```
$ gpg --detach-sign --local-user user_key -o run.sh.sig run.sh
```

Upload and submit your Job: `$ scp run.sh run.sh.sig <hpc-uid>@gwdu101.gwdg.de:/scratch/users/<hpc-uid>/home/`

```
$ ssh <hpc-uid>@gwdu101.gwdg.de 'cd /scratch/users/<hpc-uid>/home ; /opt/slurm/bin/sbatch -p secure --exclusive run.sh'
```

You can also check your Slurm output file in `/scratch/users/<hpc-uid>/home` to see if it worked.

Optional: Getting to Know the Job 8: Tutorial (10 min)

The command that we actually want to execute is

```
$ ./bart_fft.sif ../data/shepp_logan
```

which will do `bart fft -u -i`, i.e. an inverse Fourier Trafo. You can look at the images using `bart toimg`. At first, you have an frequency/k-space image:




Figure 2: K-Space image of Shepp-Logan



Figure 3: K-Space image of Shepp-Logan